

Mehrdad Fazli

Ph.D. Candidate in Computer Science, George Mason University

✉ mehردادfazli.2011@gmail.com · 📍 Falls Church, VA · 🌐 Website

🌐 LinkedIn · 🐙 GitHub · 🏠 Google Scholar

SUMMARY

Curious and rigorous researcher/engineer who digs into why models fail and turns that into concrete fixes. Strong at bridging research and production systems, whether that means diagnosing live infrastructure failures or making LLM outputs more faithful under real latency constraints.

EDUCATION

George Mason University

Ph.D. in Computer Science · Advisor: Prof. Ziwei Zhu

Fairfax, VA

Expected Aug 2027

Research: hallucination and faithfulness in large vision-language models; evaluation; explainable LLMs.

University of Virginia

M.Sc. in Systems Engineering

Charlottesville, VA

May 2023

EXPERIENCE

Google

Ph.D. Software Engineering Intern, Google Cloud Platform (Persistent Disk)

Kirkland, WA

May 2026 – Aug 2026

- Designed and implemented a multi-agent AI framework to automate slow-IO debugging and root-cause analysis for GCP's backend storage, saving hundreds of engineering hours per year.
- Triaged 20 production incidents over 3 weeks with the framework.

Amazon

Applied Scientist Intern, Personalization (P13N), Product Intent

New York, NY

Dec 2025 – Feb 2026

- Fine-tuned LLaMA 3.2 3B and Qwen3 4B (LoRA) with Claude Sonnet 4.5 as the teacher to improve the faithfulness and personalization of generated explanations for recommended items.
- Trained the students with both supervised fine-tuning (SFT) and reinforcement learning (RL); the teacher-student distillation framework let the student models match teacher performance at 3× lower latency.

Wayfair

Data Science and Machine Learning Intern, Margin Optimization

Boston, MA

Jun 2022 – Aug 2022

- Added confidence intervals to the demand simulator's short-term revenue and profit estimates to quantify uncertainty.
- Built an evaluation framework on 50M+ transactions to retrospectively assess the simulator's accuracy (SQL, Python).
- Created an error-decomposition framework to quantify the contribution of different error sources (Python).

University of Virginia

Researcher and Teaching Assistant

Charlottesville, VA

Aug 2019 – May 2023

- Built deep probabilistic forecasting models (DeepAR, DeepTCN, TFT) showing wastewater viral load improves COVID-19 forecasting.
- Fit partially observed Markov process models to agent-based simulations calibrated to real COVID-19 prevalence data.
- Built a hierarchical neural network to recognize activities from smartphone/smartwatch sensor data (IHCI 2020 Best Paper).
- Modeled product diffusion over networks to find optimal prices under local network effects.
- Teaching assistant for data and information engineering, differential equations, probability, and statistics.

SKILLS

Research areas: Multimodal LLMs, vision-language models, hallucination mitigation, evaluation & benchmarking, explainable AI

Post-training: Supervised fine-tuning (SFT), reinforcement learning, LoRA / PEFT, knowledge distillation

Agentic AI: Multi-agent systems, tool use, LLM-driven debugging and root-cause analysis

ML frameworks: PyTorch, Hugging Face Transformers, scikit-learn, TensorFlow, Keras, NumPy, Pandas, SciPy

Languages & cloud: Python, SQL, C++, R, MATLAB, Bash; GCP (Vertex AI, BigQuery), AWS (SageMaker, EC2, RDS), PySpark, Git

Visualization: Matplotlib, ggplot2, R Shiny, Power BI, Tableau

PUBLICATIONS

Under review Does Playing it Safe Count as Faithfulness? Reassessing LVLM Hallucination Mitigation Methods. **M. Fazli**, S. Mansouri, M. Marvania, Z. Zhu. arXiv:2609.01888, 2026.

ACL 2026 Findings Inject to Heal: Alleviating Hallucination in LVLMs via Context Embedding Injection (CEI). **M. Fazli**, B. Wei, Z. Zhu.

AAAI 2026 • Oral Making Sense of LLM Decisions: A Prototype-based Framework for Explainable Classification (ProtoSurE). B. Wei, **M. Fazli**, Z. Zhu.

WACV 2026 CAAC: Confidence-Aware Attention Calibration to Reduce Hallucinations in Large Vision-Language Models. **M. Fazli**, B. Wei, A. Sari, Z. Zhu.

EMNLP 2026 Findings VeriSim: A Configurable Framework for Stress-Testing Medical AI Under Patient Communication Noise. S. Mansouri, M. Marvania, V. A. Shihorkar, H. N. Tran, K. Shafiei, **M. Fazli**, Y. Li, Z. Zhu.

IEEE BigData 2023 Leveraging Deep Learning to Improve COVID-19 Forecasting Using Wastewater Viral Load. **M. Fazli**, H. Shakeri.

EJOR 2023 Competitive Pricing Under Local Network Effects. European Journal of Operational Research. A. Tavasoli, **M. Fazli**, E. Ardjmand, W. A. Young II, H. Shakeri.

FUSION 2023 Exercise and Sedentary Activity Recognition Using Late Fusion: Building Adaptable Uncertain Models. E. Juarez Garcia, V. R. Rodrigues, **M. Fazli**, L. E. Barnes, N. J. Napoli.

IEEE BigData 2021 Wastewater-Based Epidemiological Modeling for Continuous Surveillance of COVID-19 Outbreak. **M. Fazli**, S. Sklar, M. D. Porter, B. A. French, H. Shakeri.

IHCI 2020 • Best Paper HHAR-net: Hierarchical Human Activity Recognition using Neural Networks. Springer LNCS. **M. Fazli**, K. Kowsari, E. Gharavi, L. Barnes, A. Doryab.

AWARDS & SERVICE

Best Paper Award , 12th International Conference on Intelligent Human-Computer Interaction (IHCI)	2020
Outstanding Graduate Teaching Assistant Award , Department of Computer Science, George Mason University	2024
Reviewer , EMNLP 2026, NeurIPS 2026	2026