

MEHRDAD FAZLI

CONTACT

Address : Falls Church, Virginia

Email :

mehrdadfazli.2011@gmail.com

WEBSITES, PROFILES

<https://www.linkedin.com/in/mehrdad-fazli-66511479/>

<https://github.com/mehrdadfazli>

<https://mehrdadfazli.github.io/>

EDUCATION

Ph.D. in Computer Science,

Expected in 08/2027

George Mason University -

Fairfax, VA

M.Sc. in Systems Engineering,

05/2023

University of Virginia -

Charlottesville, VA

SKILLS

Languages & cloud: Python, SQL,

C++, Bash · GCP (Vertex AI,

BigQuery), AWS (SageMaker),

PySpark

Post-training: Supervised fine-

tuning (SFT), LoRA / PEFT,

knowledge distillation,

reinforcement learning

Agentic AI: Multi-agent systems,

tool use, LLM-driven debugging &

root-cause analysis

Frameworks: PyTorch, Hugging

Face Transformers, scikit-learn

PROFESSIONAL SUMMARY

Forward-thinking, Curious, and rigorous researcher/engineer who digs into why models fail and turns that into concrete fixes. Strong at bridging research and production systems, whether that means diagnosing live infrastructure failures or making LLM outputs more faithful under real latency constraints.

WORK HISTORY

PhD SWE Intern, 5/2026 to 08/2026 – Kirkland, WA

 (Google Cloud Platform)

- Designed and implemented a multi-agent AI framework to automate debugging and root-cause analysis for GCP's backend storage, saving hundreds of engineering hours per year.
- Triaged 20 production incidents over 3 weeks with the framework.

Applied Scientist Intern, 12/2025 to 02/2026 - New York, NY

 (Personalization P13N – Product Intent Team)

- Finetuned LLaMA 3.2 3B and Qwen3 4B (LoRa) using Claude Sonnet 4.5 as the teacher to improve faithfulness and personalization of the generated explanations for the recommended items. The teacher-student distillation framework allowed the student models to achieve a performance on par with the teacher model at a 3X lower latency.

Data Science and Machine Learning Intern, 06/2022 to

08/2022 – Boston, MA

 **Wayfair** (Margin Optimization Team)

- Added confidence intervals to the "demand simulator" estimations of short-term revenue and profit to quantify the uncertainty in the estimations of outcomes (Python)
- Developed an evaluation framework for the "demand simulator" using over 50 million transactions to retrospectively evaluate the accuracy of the estimated outcomes (SQL and Python)
- Created a framework for error decomposition to quantify the contribution of various sources of errors (Python)

CERTIFICATIONS

- Machine Learning - Coursera
- Neural Networks and Deep Learning - Coursera
- Convolutional Neural Networks - Coursera
- Essential Design Principles for Tableau - Coursera

Researcher, 11/2021 to 05/2023 – Charlottesville, VA University of Virginia

- Fitted a partially observed Markov process (POMP, aka HMM) to a synthetic dataset generated via agent-based modeling (ABM) adjusted for the actual COVID-19 prevalence data
- Trained three different probabilistic forecasting neural network architectures, DeepAR, DeepTCN, and TFT, on COVID-19 data and performed feature analysis and data visualization to find the most important features for COVID-19 forecasting
- Demonstrated that viral load guides the forecasting models and improves their accuracy
- Price Optimization Under Local Network Effects: Stochastic modeling of the diffusion of a new product in a network to find the optimal prices

Teaching and Research Assistant, 08/2019 to 05/2021 – Charlottesville, VA University of Virginia

- Trained a deep neural network model to classify a user's activity based on the sensor data from their smartphone and smartwatches
- Assisting undergraduate students in "data and information engineering" and math courses, including "differential equations," "probability," and "statistics" via holding office hours and helping them with the assignments

SELECTED PUBLICATIONS

- Fazli, Mehrdad, Bowen Wei, and Ziwei Zhu. "[CAAC: Confidence-Aware Attention Calibration to Reduce Hallucinations in Large Vision-Language Models](#)" **WACV 2026**
- Wei, Bowen, Mehrdad Fazli, and Ziwei Zhu. "[Making Sense of LLM Decisions: A Prototype-based Framework for Explainable Classification](#)" **AAAI 2026**
- Fazli, Mehrdad, Bowen Wei, and Ziwei Zhu. " [Inject to Heal: Alleviating hallucination in LVLMS via Context Embedding Injection](#)" **ACL 2026**
- Tavasoli, Ali, Mehrdad Fazli, et. al. "[Competitive Pricing Under Local Network Effects](#)" **European Journal of Operational Research (2023)**.
- Fazli, Mehrdad, and Heman Shakeri. " [Leveraging Deep Learning to Improve COVID-19 Forecasting Using Wastewater Viral Load](#)" **IEEE Big Data 2023**.
- Fazli, Mehrdad, et. al. " [Hallucination Mitigation or Conservative Decoding? Rethinking Faithfulness Evaluation in LVLMS.](#)" **[under review for EMNLP 2026]**.

AWARDS AND HONORS

- **Outstanding graduate teaching assistant** (GTA) award – computer science department GMU - 2024
- **Best paper** award at 12th International Conference of Intelligent Human Computer Interaction - 2021